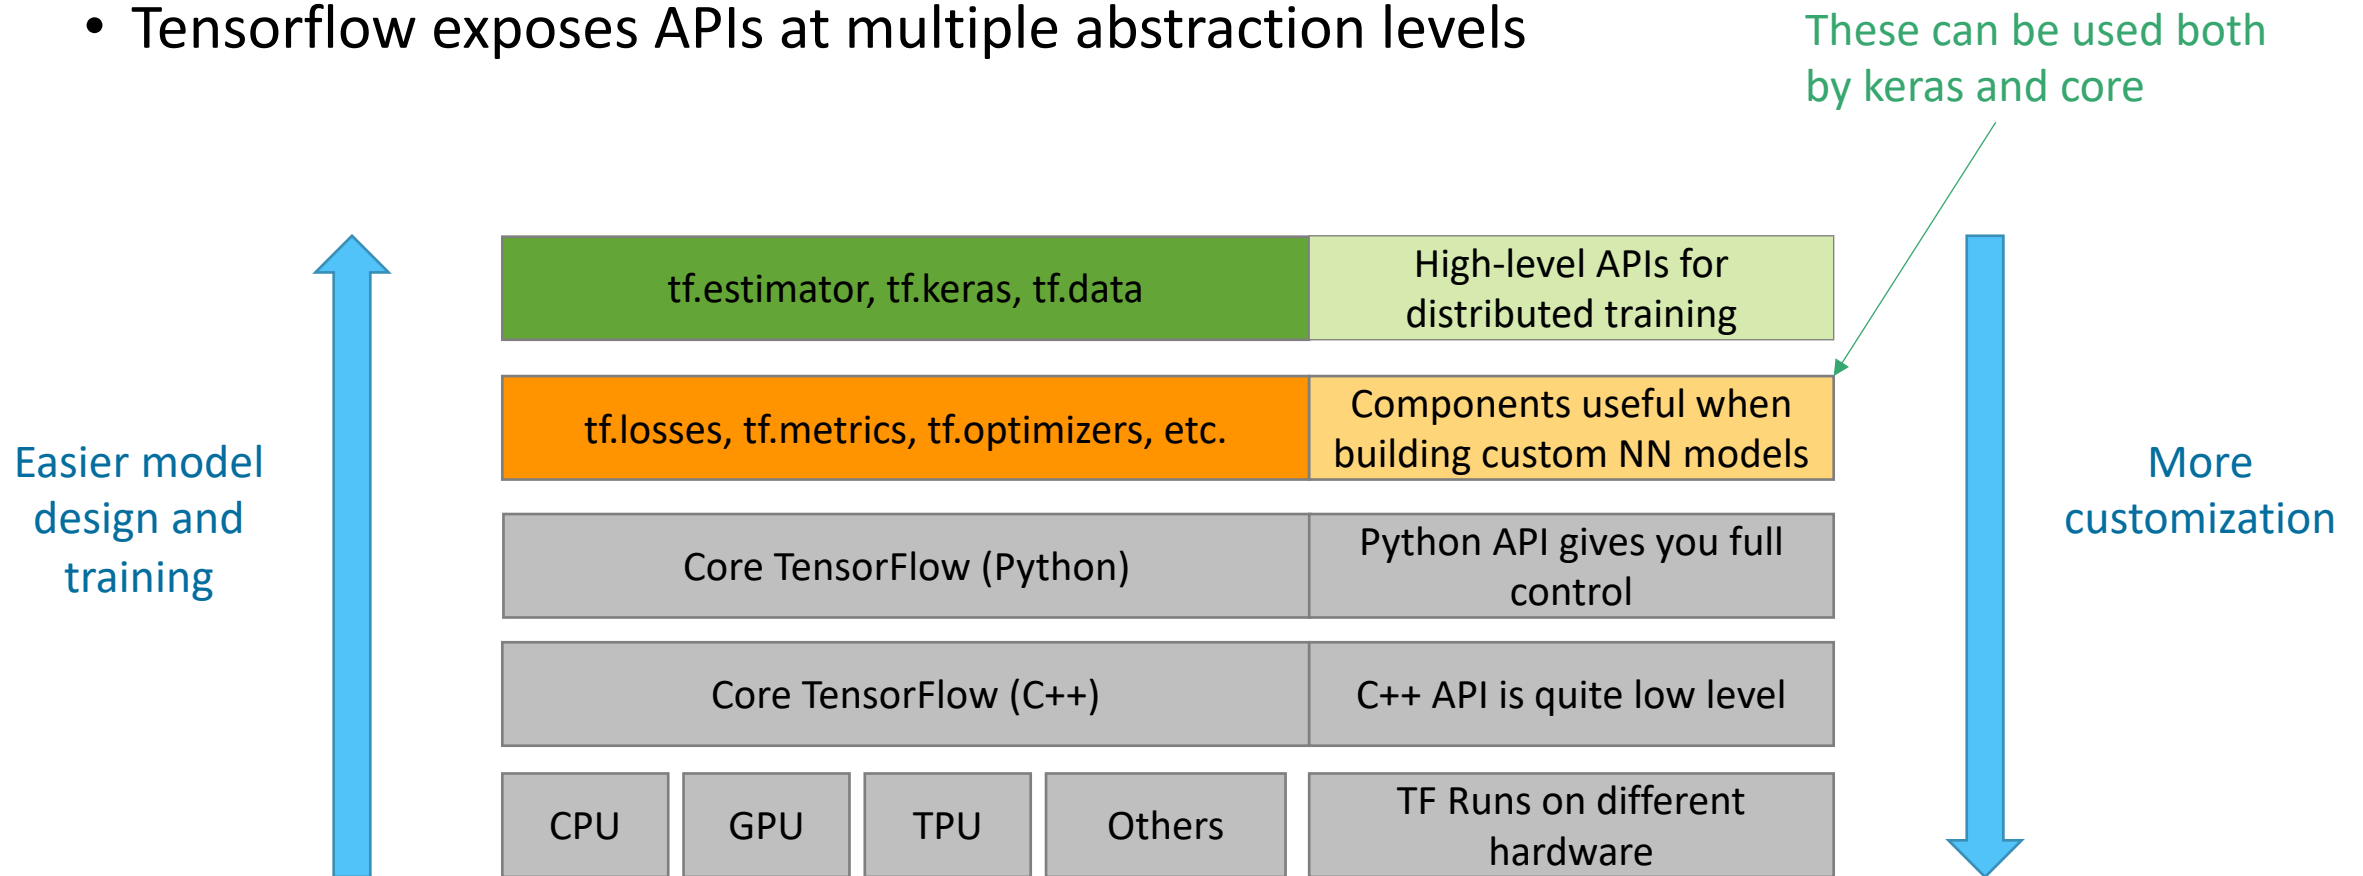


Tensorflow 2

Keras API (Continued)

Tensorflow API Hierarchy

- Tensorflow exposes APIs at multiple abstraction levels




The Keras API


- Most of the times you don't need to write custom models and training loops from scratch
- Keras is a high-level API built-in in TF2, that provides a much easier flow to write “standard” models
 - No need to worry about gradient tapes, weight updates, etc. manually.
- A Keras model can be built in 3 main ways:
 - **Sequential API** → for Single-Input Single-Output models built as stacks of layers
 - **Functional API** → for MIMO models, residual connections, etc.
 - **Sub-classing API** → for maximum customization (e.g. dynamic/adaptive models)

The Keras API

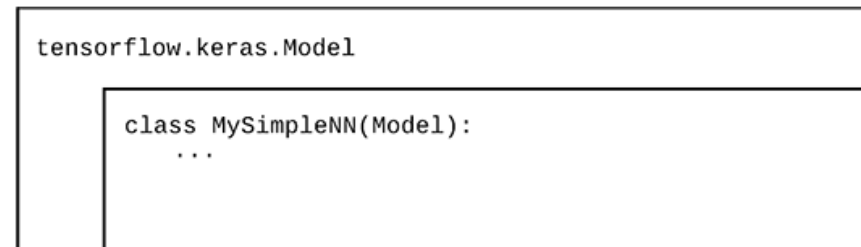
1. Sequential API



2. Functional API



3. Model Subclassing

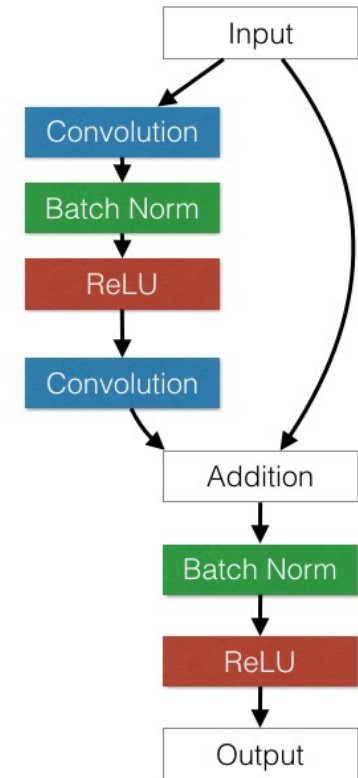
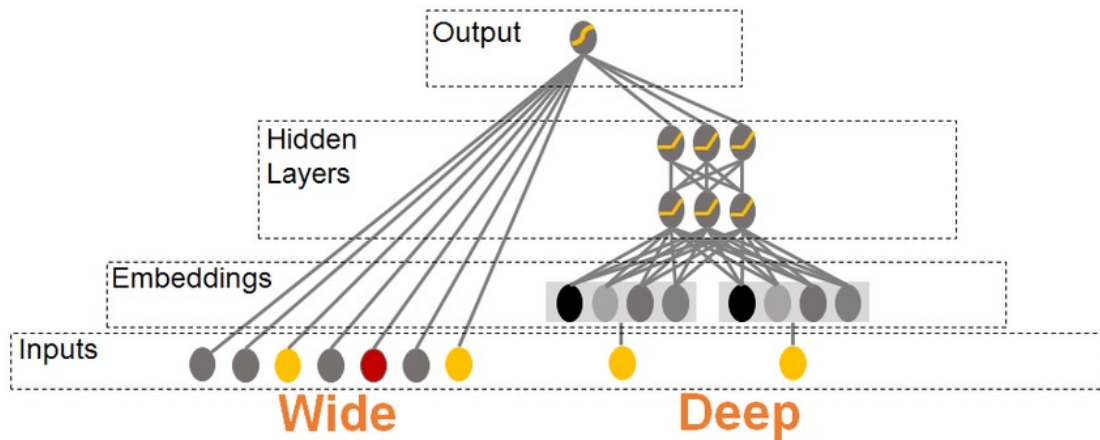


Tensorflow 2

Keras Functional API

The Keras Functional API

- The Sequential API is limited to single-in single-out models built as a stack of layers.
 - The most common scenario, but not the only one.
 - E.g. deep and wide models, residual connections, shared layers, etc:



The Keras Functional API

- The Functional API can handle arbitrary (static) graphs of layers
- First create an input node (implicit in the Sequential API):

```
inputs = keras.Input(shape=(28, 28))
```
- Then create new layer instances and apply them to the input by calling them as functions:

```
flatten = keras.Flatten(input_shape=(28, 28))  
x = flatten(inputs)
```
- Or more synthetically:

```
x = keras.Flatten(input_shape=(28, 28))(inputs)
```

The Keras Functional API

- Connect many of these layer “nodes” as you wish:

```
inputs = keras.Input(shape=(28, 28))
x = keras.Flatten(input_shape=(28, 28))(inputs)
x = layers.Dense(64, activation="relu")(x)
x = layers.Dense(64, activation="relu")(x)
outputs = layers.Dense(10)(x)
```

- Finally create a `Model()` by specifying the input(s) and output(s) of the graph:

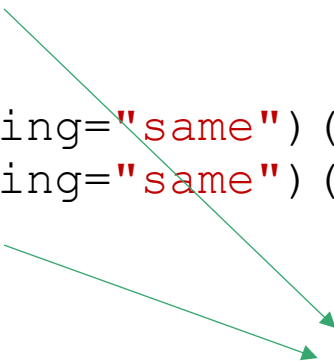
```
model = keras.Model(
    inputs=inputs,
    outputs=outputs,
    name="my_model")
```


The Keras Functional API

```
inputs = keras.Input(shape=(32, 32, 3), name="img")
x = layers.Conv2D(32, 3, activation="relu")(inputs)
x = layers.Conv2D(64, 3, activation="relu")(x)
block_1_output = layers.MaxPooling2D(3)(x)

x = layers.Conv2D(64, 3, activation="relu", padding="same")(block_1_output)
x = layers.Conv2D(64, 3, activation="relu", padding="same")(x)
block_2_output = layers.add([x, block_1_output])

x = layers.Conv2D(64, 3, activation="relu", padding="same")(block_2_output)
x = layers.Conv2D(64, 3, activation="relu", padding="same")(x)
block_3_output = layers.add([x, block_2_output])
```



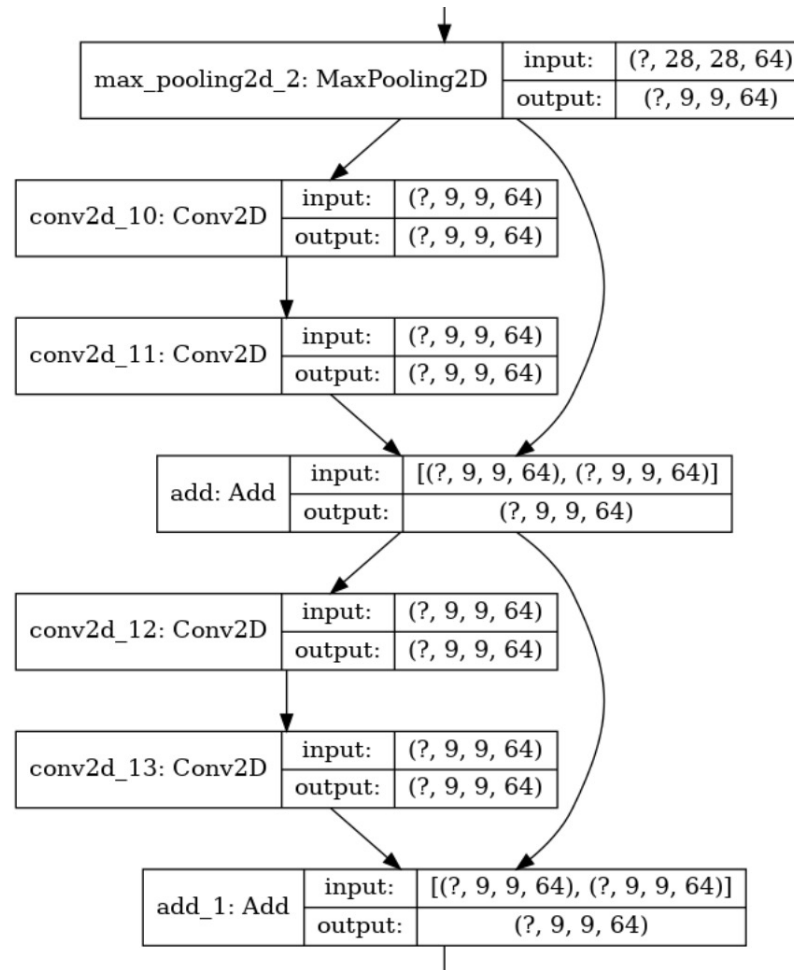
Add input and output to
create a residual connection

The Keras Functional API

```
x = layers.Conv2D(64, 3, activation="relu")(block_3_output)
x = layers.GlobalAveragePooling2D()(x)
x = layers.Dense(256, activation="relu")(x)
x = layers.Dropout(0.5)(x)
outputs = layers.Dense(10)(x)

model = keras.Model(inputs, outputs, name="toy_resnet")
```

The Keras Functional API



The Keras Functional API

- Only step 2 of the whole Keras flow changes, the rest is identical!
- **Notebook:** Keras Functional API

Tensorflow 2

Keras Sub-Classing API

The Keras Sub-classing API

- Useful when you want to build models that:
 - Don't use standard layers provided by keras (e.g. Conv, MaxPooling, GRU, LSTM, BatchNormalization, Dropout, etc.)
 - Or cannot be represented by DAGs (e.g. tree RNNs)
 - We won't go in so many details, but we'll show you the API for completeness (and because it's similar to PyTorch)
- Custom layers are created by sub-classing the `keras.Layer` class
- Custom models are created by sub-classing the `keras.Model` class

The Keras Sub-classing API

- **CAUTION:** in general, don't use the sub-classing API unless you really need it
- Sequential/Functional models are much easier to:
 - Inspect (`summary()`, `plot_model()`, etc.)
 - Debug (input shape and dtype specified in advance using `keras.Input`).
 - Serialize/clone/save/restore (with subclassing, you have to write your own model saving code)



The Keras Sub-classing API

- A custom layer = some trainable `tf.Variable` and some computations:

```
class MyLinear(keras.layers.Layer):
```

Base class
constructor

```
    def __init__(self, units=32):  
        super(MyLinear, self).__init__()  
        self.units = units
```

```
    ...
```


The Keras Sub-classing API

- A custom layer = some trainable `tf.Variable` and some computations:

```
class MyLinear(keras.layers.Layer):
```

```
...
```

```
def build(self, input_shape):
```

```
    w_init = tf.random_normal_initializer()
```

```
    self.w = tf.Variable(
        initial_value=w_init(shape=(input_shape[-1], self.units),
                               dtype="float32"), trainable=True)
```

```
    b_init = tf.zeros_initializer()
```

```
    self.b = tf.Variable(
        initial_value=b_init(shape=(self.units,),
                               dtype="float32"), trainable=True)
```

```
...
```

Called only the first time a layer is invoked.
Useful to support variable input shapes

Define and
initialize weights

Define and
initialize biases

The Keras Sub-classing API

- A custom layer = some trainable `tf.Variable` and some computations:

```
class MyLinear(keras.layers.Layer):
```

```
...
```

```
def call(self, inputs):  
    return tf.matmul(inputs, self.w) + self.b
```



A set of TF ops for
autograd

The Keras Sub-classing API

- You can then use this custom layer in a standard (functional or sequential) way:

```
x = tf.ones((2, 2))  
linear_layer = MyLinear(4)  
y = linear_layer(x)
```

The Keras Sub-classing API

- You can avoid the explicit creation of `tf.Variable` thanks to the `add_weight()` method:

```
def build(self, input_shape):  
    super(Linear, self).__init__()  
  
    self.w = self.add_weight(shape=(input_shape[-1], self.units),  
                             initializer="random_normal", trainable=True)  
  
    self.b = self.add_weight(shape=(self.units,),  
                             initializer="zeros", trainable=True)
```

- You can set `trainable=False` if you want non-trainable weights

The Keras Sub-classing API


- Custom layers can then be used as part of other custom layers (recursively composable).
- You can define two other functions (`get_config()` and `from_config()`) to support serialization.
- You can track custom losses and metrics, and many other advanced functionalities.
- See full details here:
https://www.tensorflow.org/guide/keras/custom_layers_and_models

The Keras Sub-classing API

- Sub-classes of `keras.Models` define entire models (i.e. objects you will train)
- Should I use the Layer class or the Model class?
 - will I need to call `fit()` on it?
 - Will I need to call `save()` on it?
 - If so, go with Model.

The Keras Sub-Classing API

- Example (very similar to PyTorch):

```
class ResNet(tf.keras.Model):  
  
    def __init__(self):  
        super(ResNet, self).__init__()  
        self.block_1 = ResNetBlock()   
        self.block_2 = ResNetBlock()  
        self.global_pool = layers.GlobalAveragePooling2D()  
        self.classifier = Dense(num_classes)  
  
    def call(self, inputs):  
        x = self.block_1(inputs)  
        x = self.block_2(x)  
        x = self.global_pool(x)  
        return self.classifier(x)
```

The Keras Sub-Classing API

- **Notebook:** Keras Subclassing API

Tensorflow 2

(Some) Advanced Keras Features

(Some) Advanced Keras Features

1. Saving Models
2. Advanced options for `compile()`:
 - Custom metrics and losses
3. Advanced options for `fit()`:
 - Creating/passing validation sets
 - Class and sample weights
 - Using callbacks
4. Tensorboard

(Some) Advanced Keras Features

1. Saving Models
2. Advanced options for `compile()`:
 - Custom metrics and losses
3. Advanced options for `fit()`:
 - Creating/passing validation sets
 - Class and sample weights
 - Using callbacks
4. Tensorboard

Saving Models

- A Keras model consists of multiple components:
 - An **architecture**, or configuration, which specifies what layers the model contain, and how they're connected.
 - A set of **weights** values (the "state of the model").
 - An **optimizer** (defined by compiling the model).
 - A set of **losses and metrics** (defined by compiling the model)
- You can save all these components or only some of them.

Saving Models

- Saving everything:

```
model.save('path/to/location')  
  
# load back  
model = keras.models.load_model('path/to/location')
```

- The reconstructed model is already compiled and has retained the optimizer state. You can safely resume training.

Saving Models

- Saving only the architecture (no weights, no compile information):

```
json_config = model.to_json()  
new_model = keras.models.model_from_json(json_config)
```

- The reconstructed model weights are re-initialized and there is no optimizer info, etc.

Saving Models

- Saving only the weights (e.g. for checkpointing):

```
model.save_weights("ckpt")  
load_status = model.load_weights("ckpt")
```

- By default uses TF Checkpoint format.

(Some) Advanced Keras Features

1. Saving Models
2. Advanced options for `compile()`:
 - Custom metrics and losses
3. Advanced options for `fit()`:
 - Creating/passing validation sets
 - Class and sample weights
 - Using callbacks
4. Tensorboard

Custom Losses and Metrics

- You can pass the `compile()` function a custom loss function:

```
def my_mse(y_true, y_pred):  
    return tf.reduce_mean(tf.square(y_pred - y_true))
```

```
model.compile(optimizer="adam", loss=my_mse, metrics=...)
```

Custom Losses and Metrics

- You can pass the `compile()` function one or more custom metrics, and combine them with built-in ones:

```
model.compile(optimizer="adam", loss=..., metrics=[my_mse, "mae"])
```

Custom Losses and Metrics

- **Notebook:** Regression with Custom Loss & Metrics

(Some) Advanced Keras Features

1. Saving Models
2. Advanced options for `compile()`:
 - Custom metrics and losses
3. Advanced options for `fit()`:
 - Creating/passing validation sets
 - Class and sample weights
 - Using callbacks
4. Tensorboard

Creating Validation Sets

- Validation data can be either passed to the `fit()` function....

```
history = model.fit(  
    X_train,  
    Y_train,  
    epochs=10,  
    batch_size=32,  
    validation_data=<validation_data>,  
)
```

- The main accepted formats are: a tuple of (x_val, y_val) or a TF dataset

Creating Validation Sets

- ...or we can let fit() pick a validation set for us:

```
history = model.fit(  
    X_train,  
    Y_train,  
    epochs=10,  
    batch_size=32,  
    validation_split=<fraction in [0-1], e.g. 0.2>,  
)
```

- In this case, a fraction of the training data will be used for validation (and the model won't be trained on it)

Creating Validation Sets

- By default, validation is performed after every epoch:

```
Epoch 1/40
8/8 [=====] - 0s 14ms/step - loss: 582.3579 - mae: 22.8670 - val_loss: 581.9780 - val_mae: 2
2.9370
Epoch 2/40
8/8 [=====] - 0s 3ms/step - loss: 539.1675 - mae: 22.0812 - val_loss: 537.9955 - val_mae: 2
2.1408
```

- You can change this with the parameter `validation_freq`

Class and Sample Weights

- You can pass a dictionary of: `{class_index : class_weight}` to `fit()`, in order to weigh the loss function based on the true class of a sample.
- As you probably know, this is one of the ways to deal with class imbalance, letting the training algorithm give “more importance” to under-represented samples.

```
history = model.fit(  
    X_train,  
    Y_train,  
    epochs=10,  
    batch_size=32,  
    class_weight={0: 1., 1: 50., 2: 2.}  
)
```


Class and Sample Weights

- Alternatively, you can provide fit() with a weight for every training sample, which is useful (for example) when you know that some samples' labels have higher “confidence” than others
- The basic way to pass sample weights is using an array of the same length as the training data

```
history = model.fit(  
    X_train,  
    Y_train,  
    epochs=10,  
    batch_size=32,  
    sample_weight=<tensor or np.array>)
```

Callbacks


- Callbacks are a powerful tool to customize the behavior of Keras during training, evaluation and inference.
- Callbacks can be passed to `fit()`, `evaluate()` and `predict()`
- All callbacks are sub-classes of `keras.callbacks.Callback` and override a set of methods called at various stages of the three functions above

Callbacks

- Keras provides you with a large set of pre-cooked callbacks, such as:
 - `EarlyStopping`: Stop training when a monitored metric has stopped improving
 - `ModelCheckpoint`: Save the Keras model or model weights with some frequency
 - `LearningRateScheduler`: change the LR during training
 - `TensorBoard`: Enable visualizations for TensorBoard (see later)
 - Many others...

Callbacks (Early Stopping)

- Early stopping is a good practice for training deep learning models, useful to:
 - Avoid computationally intensive training epochs when they do not improve the predictive capabilities of the model.
 - Stop training in the so-called “optimal capacity point, where bias and variance are both low.



Callbacks (Early Stopping)

- Some parameters:

```
tf.keras.callbacks.EarlyStopping(  
    monitor='val_loss',    Quantity to monitor  
    patience=0,           Number of not-improving epochs before stopping  
    verbose=0,           Print a message explaining why you stopped  
    mode='auto',         Specify if the target quantity should be minimized/maximized  
    ...                  (auto uses the name)
```

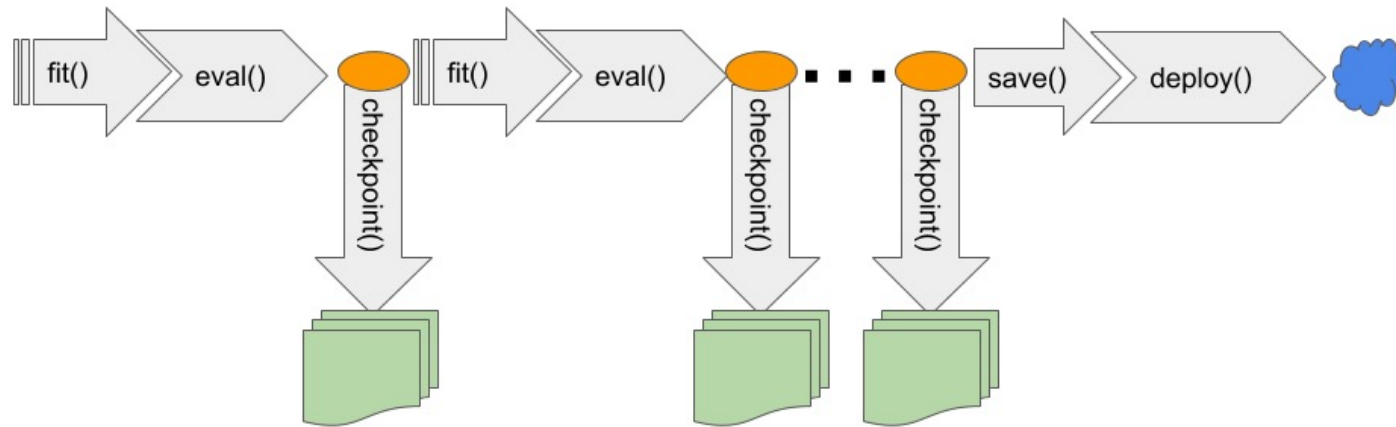
Callbacks (Early Stopping)

- Example of usage:

```
history = model.fit(  
    X, Y, epochs=200, ...,  
    callbacks=[EarlyStopping(  
        monitor='val_loss',  
        patience=2,  
        verbose=1)  
    ])
```

Callbacks (Model Checkpoint)

- Checkpoints are simply periodic snapshots of the current status of a model, saved during training. Useful to:
 - Avoid losing everything in case of crashes
 - Analyzing the model behavior over the training epochs
 - Retrieving the overall best model for deployment



Callbacks (Model Checkpoint)

- Some parameters:

```
tf.keras.callbacks.ModelCheckpoint(  
    filepath, Destination  
    monitor='val_loss', Metric to monitor (e.g. for save_best_only)  
    verbose=0,  
    save_best_only=False, Don't overwrite best checkpoint version  
    save_weights_only=False, Save with model.save() or with model.save_weights()  
    mode='auto', Same as EarlyStopping  
    save_freq='epoch', 'epoch' or an integer (interpreted as number of batches)  
)
```


Callbacks (Model Checkpoint)

- Example of usage:


```
cp_callback = ModelCheckpoint(  
    './callback_test_chkp/chkp_{epoch:02d}_{val_loss:.2f}',  
    save_best_only=False,  
    save_weights_only=False,  
    save_freq='epoch'  
)
```



The filepath can use named formatting options for the epoch, loss and metrics values.


Callbacks (Learning Rate Scheduling)

- The Learning Rate (LR) is one of the most important hyper-parameters of a deep learning model.



Callbacks (Learning Rate Scheduling)

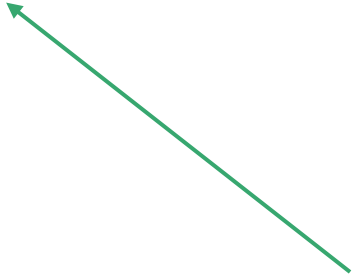
- LR Scheduling refers to dynamically changing the LR during training. There are many possible schedules proposed (still very much an art...)
- Example: reduce LR on plateau



Callbacks (Learning Rate Scheduling)

- Some parameters:

```
lr_callback = LearningRateScheduler(  
    schedule,  
    verbose=0,  
)
```



A Python function that takes the epoch and the current LR and returns the next LR

Callbacks (Learning Rate Scheduling)

- Example of usage:

```
def my_schedule(epoch, lr):  
    if epoch < 10:  
        return lr  
    else:  
        return lr * tf.math.exp(-0.1)
```

```
lr_callback = LearningRateScheduler(my_schedule, verbose=0)
```

Callbacks (Learning Rate Scheduling)

- There's also a pre-cooked LR callback for “Reduce on plateau” technique:


```
reduce_lr = tf.keras.callbacks.ReduceLROnPlateau(  
    monitor='val_loss',  
    factor=0.3,  
    patience=5,  
    min_lr=1e-04,  
    verbose=1  
)
```

Callbacks (Custom)

- You can create your own callbacks by subclassing the `Callback` class.
- Methods that can be overridden:
 - **`on_(train|test|predict)_begin(self, logs)`**: Called at the beginning of fit/evaluate/predict.
 - **`on_(train|test|predict)_end(self, logs)`**: Called at the end of fit/evaluate/predict.
 - **`on_(train|test|predict)_batch_begin(self, batch, logs)`**: Called right before processing a batch during training/testing/predicting.
 - **`on_(train|test|predict)_batch_end(self, batch, logs)`**: Called at the end of training/testing/predicting a batch.
 - **`on_epoch_begin(self, epoch, logs)`**: Called at the beginning of an epoch during training.
 - **`on_epoch_end(self, epoch, logs)`**: Called at the end of an epoch during training.

Callbacks (Custom)

- Example: `on_epoch_end(self, epoch, logs)`



`self.model` is a pointer to the model instance, which can be used to stop training, implement LR scheduling, etc.

epoch number

Stores metrics and losses values as a dictionary of: {quantity_name: val}, e.g. `logs['mae']` or `logs['loss']`

Callbacks (Custom)

- Example: custom early-stopping the first time the epoch accuracy reaches 90%:

```
class myEarlyStopping(tf.keras.callbacks.Callback):  
    def on_epoch_end(self, epoch, logs):  
        if(logs.get('accuracy')>0.9):  
            print("\nReached 90% accuracy so cancelling training!")  
            self.model.stop_training = True
```

Callbacks

- **Notebook:** Callbacks

Using Datasets with Keras

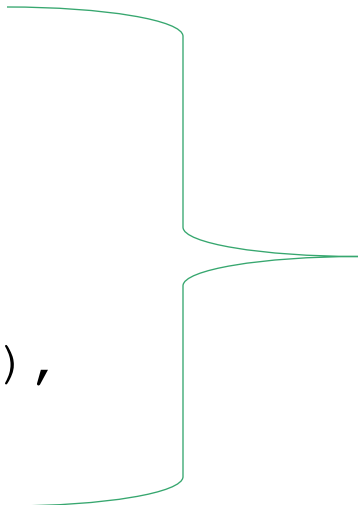
- Of course, the Keras API can digest `tf.data.Datasets` naturally.
- One important caveat: by default, you should not call `repeat()` to make an infinite dataset. Repetition is handled internally by Keras
 - If you already have an infinite dataset, you must pass the additional parameters `steps_per_epoch` parameter in `fit()` and `steps` in `evaluate()`
- Note that the `predict()` function does not need labels, but if you pass a dataset that contains two components, the labels are simply ignored.

Using Datasets with Keras

- Moreover, the `keras.preprocessing` package contains many pre-cooked input loading functions that produce Datasets, hiding all the details of `tf.data`

- Example:

```
train_ds = tf.keras.preprocessing.image_dataset_from_directory(  
    data_dir,  
    validation_split=0.2,  
    subset="training",  
    seed=123,  
    image_size=(img_height, img_width),  
    batch_size=batch_size)
```



Train/val split, batching,
resizing in a single step.

Using Datasets with Keras

- **Notebook:** TF Data and Keras