

Machine Learning for IoT

Tensorboard

(Some) Advanced Keras Features

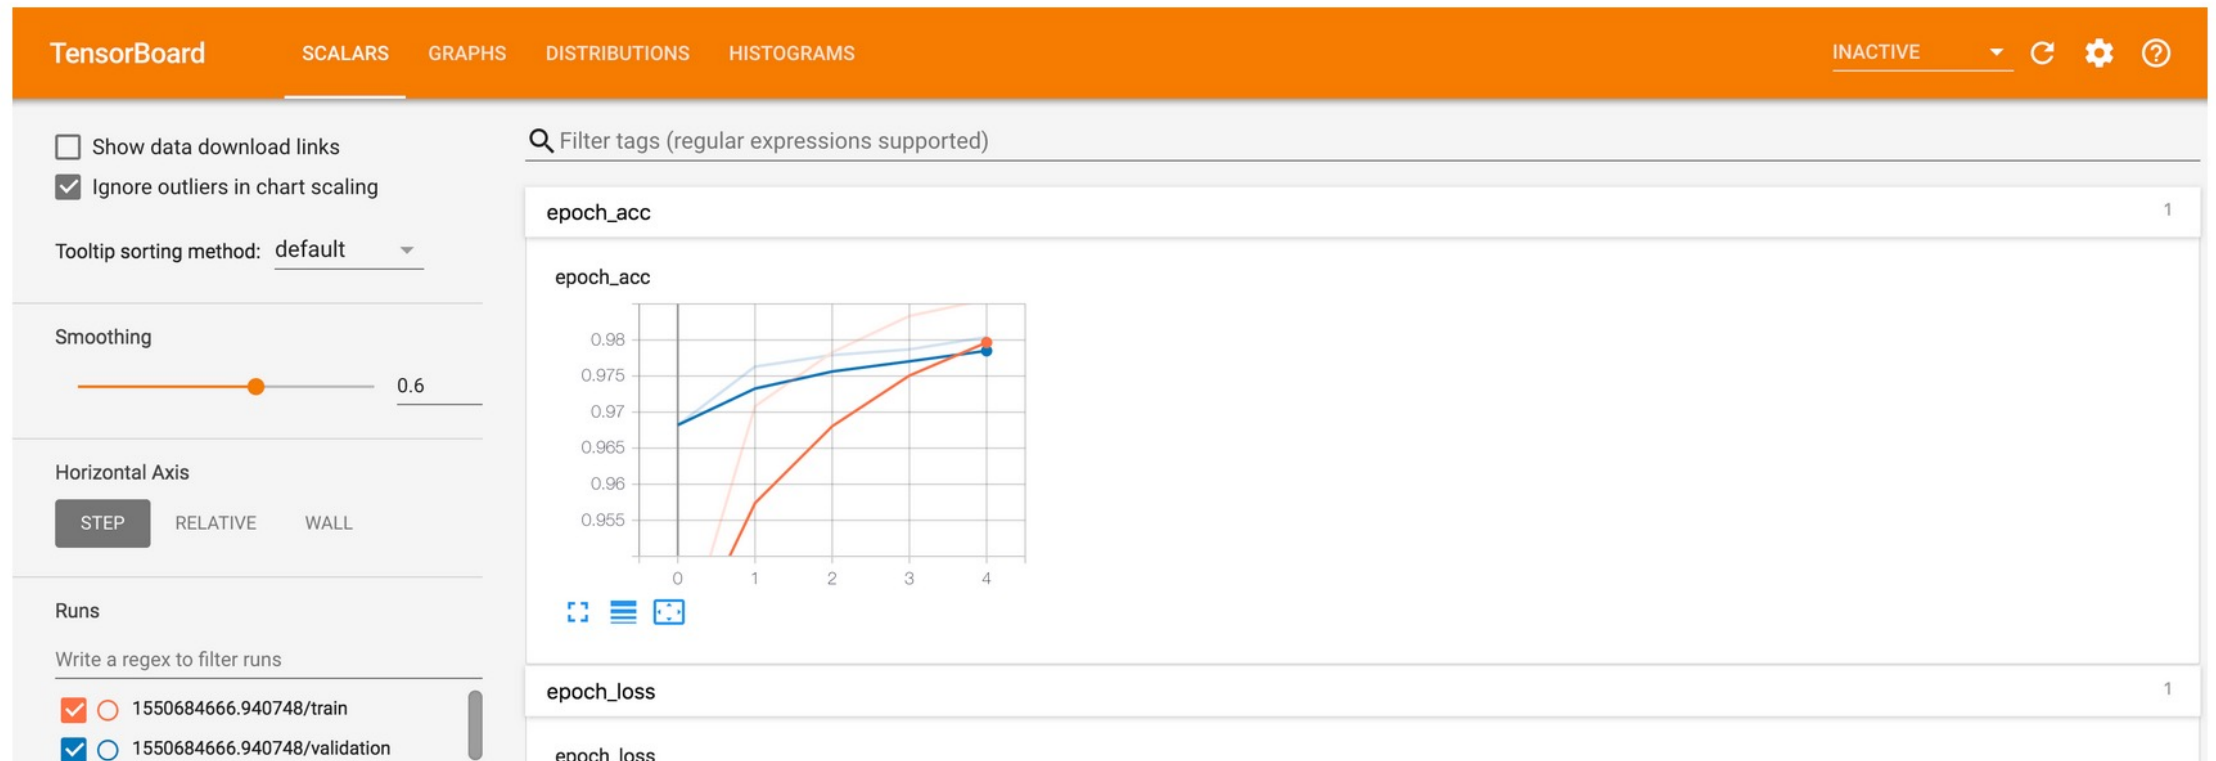
1. Saving Models
2. Advanced options for `compile()`:
 - Custom metrics and losses
3. Advanced options for `fit()`:
 - Creating/passing validation sets
 - Class and sample weights
 - Using callbacks
4. Tensorboard

TensorBoard

- TensorBoard is an analysis and visualization tool often used together with TF and Keras (but not only).
- It allows you to:
 - Track the loss and metrics during training
 - View the model computation graph
 - Etc.


TensorBoard

- The TensorBoard Dashboard



TensorBoard

- The **Scalars** dashboard track loss and metrics in every epoch. You can use it to also track training speed, learning rate, and other scalar values.
- The **Graphs** dashboard helps you visualize your model.
- The **Distributions** and **Histograms** dashboards show the distribution of a Tensor over time. This can be useful to visualize weights and biases and verify that they are changing in an expected way.



- You can also enable TensorBoard to log images, embeddings, etc.

TensorBoard

- You can enable TensorBoard logging in Python using a dedicated Callback

```
tf.keras.callbacks.TensorBoard(  
    log_dir='logs', Destination  
    histogram_freq=0, How frequently (epochs) to compute weight histograms (never if 0)  
    write_graph=True, Save the graph in TB  
    write_images=False, Save model weights as images in TB  
    update_freq='epoch', 'batch', 'epoch' or integer number of batches  
    ...  
)
```

TensorBoard

- Then after starting the training (also while it is running) you can check the results in Tensorboard by starting it from the command line:

```
tensorboard --logdir <path specified in the callback>
```

- Or from a Jupyter Notebook (same command with a % at the beginning):
- Note that it is convenient to use a single logdir with sub-directories for different training runs.

TensorBoard

- You can also manually log scalars in TB format like this:

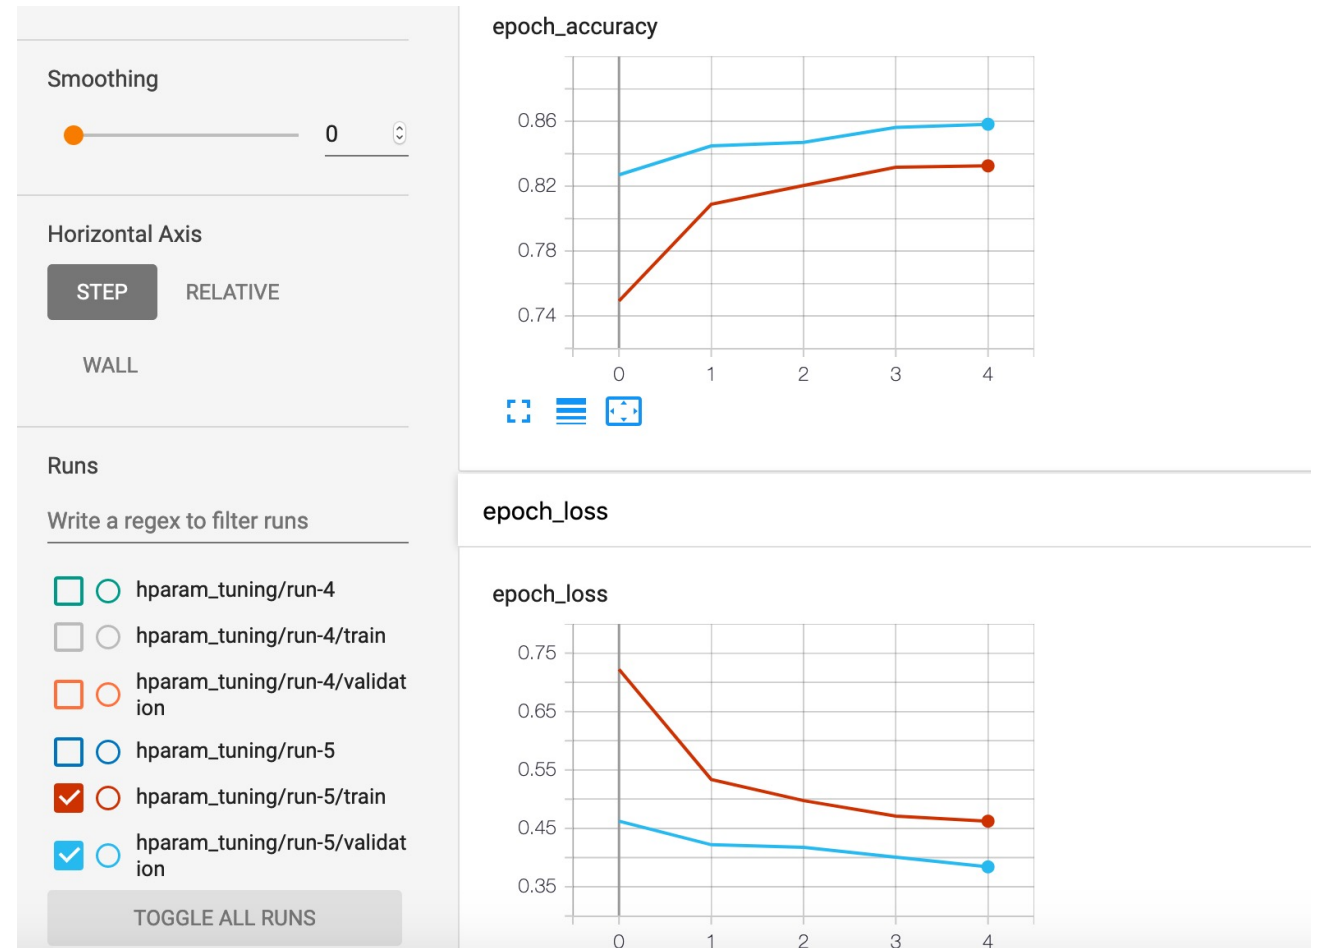
```
train_writer= tf.summary.create_file_writer(log_dir + "/train")
```

Use separate train/val/test subfolders for better visualization

```
# ... Some code (e.g. custom training loop...)
with train_summary_writer.as_default():
    tf.summary.scalar('loss', train_loss, step=epoch)
```


TensorBoard

- Tensorboard Scalar Logs



TensorBoard

- **Notebook:** Tensorboard_Scalars

TensorBoard (Hyper-parameters Search)

- TensorBoard is helpful for hyper-parameters search:
 - This is one of the most important steps in deep learning models training.
 - Although there are tools for doing this automatically using various optimization algorithms (Neural Architecture Search/AutoML), in some cases it is still done manually due to resource limitations.
- TensorBoard can visualize the results of multiple training runs with different hyper-parameters in a convenient way, so to guide designers in their next choices.

TensorBoard (Hyper-parameters Search)

1. Import hyper-parameters API:

```
from tensorboard.plugins.hparams import api as hp
```

2. Create a dictionary of hyper-parameters:

```
hparams = {  
    'num_units': num_units,  
    'dropout': dropout_rate,  
}
```

3. Create a Model() using those hparams.


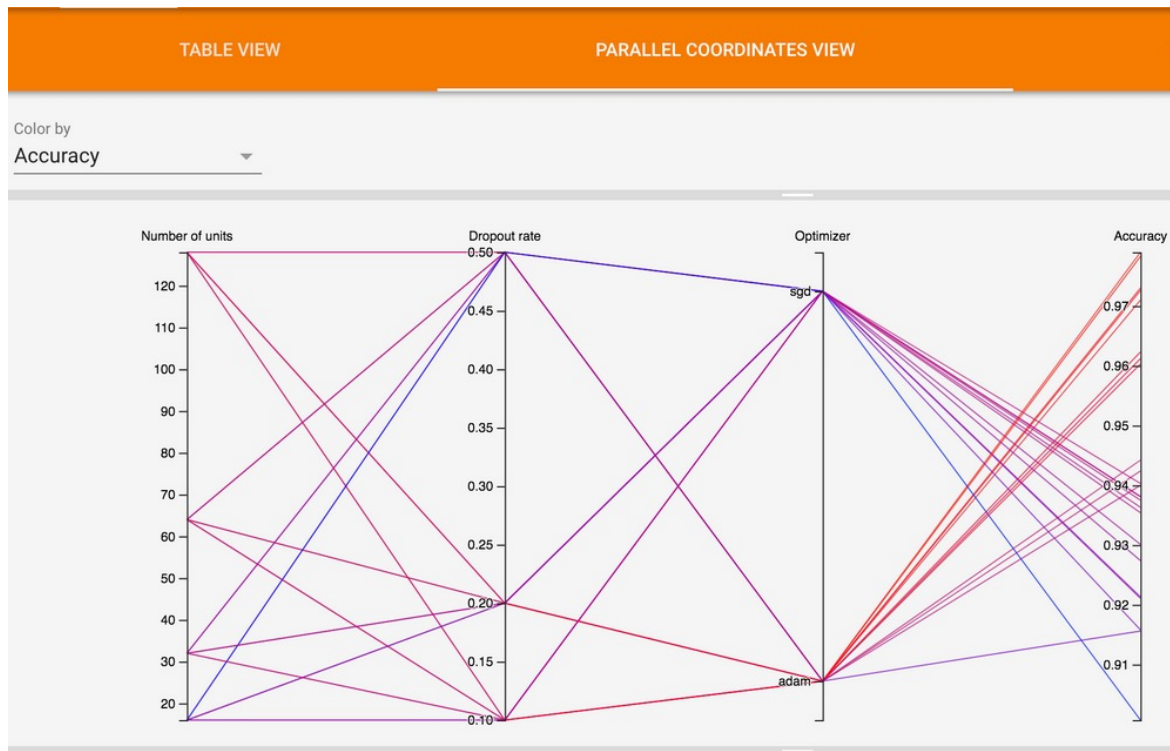
TensorBoard (Hyper-parameters Search)

4. Use a pre-defined callback to log the training metrics associated with the current hyper-parameters, and pass it to `fit()`

```
hp_callback = hp.KerasCallback(log_dir, hparams)
model.fit(
    x_train,
    y_train,
    validation_split=0.1,
    epochs=5,
    callbacks=[tb_callback, hp_callback]
)
```

TensorBoard (Hyper-parameters Search)

- See the search results in TensorBoard with different views



TensorBoard (Hyper-parameters Search)

- **Notebook:** TensorBoard_Hyperparameters

TensorBoard (Profiling)

- Another useful functionality of TensorBoard is the ability to profile the execution of a model (e.g. during training) to identify and remove performance bottlenecks.
- This is done through `tensorboard_plugin_profile` (should be pre-installed on newest TF/TB versions)

TensorBoard (Profiling)

- You can enable profiling of training steps adding the `profile_batch` option to the TensorBoard callback.

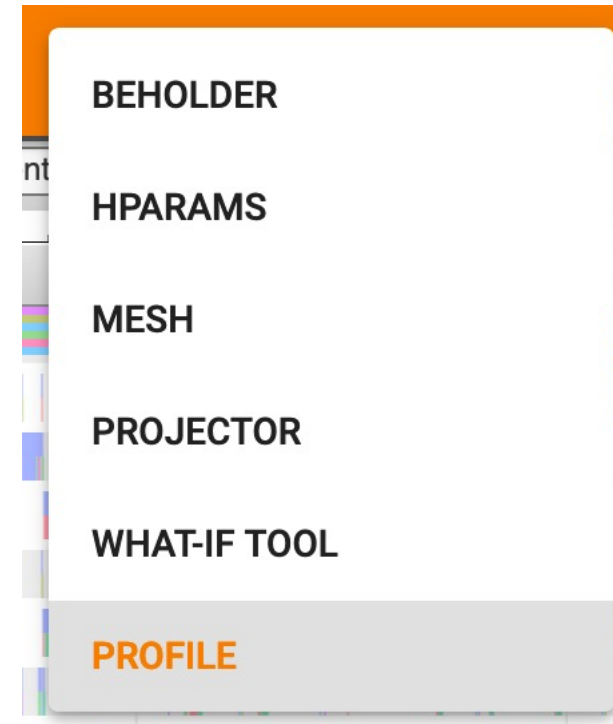
```
tb_callback = tf.keras.callbacks.TensorBoard(  
    log_dir = './tb_logs',  
    histogram_freq = 1,  
    profile_batch = '500,520'  
)
```

Profile batches from 500 to 520.




TensorBoard (Profiling)

- After training, result appear in the “Profile” tab of TensorBoard (find it in the drop-down menu if not already visible).




TensorBoard (Profiling)

- The “Overview Page”. shows a breakdown of the time taken by various categories of operations in the profiled batches (20 in this example).
 - It also provides useful textual tips on how to improve the model performance




TensorBoard (Profiling)

- The “Trace View” pane shows the trace of individual events with the corresponding execution time, grouped by type




TensorBoard (Profiling)

- The “TensorFlow Stats” pane shows summary statistics.



Healthy Profile

Troublesome Profile
(input bound)



TensorBoard (Profiling)

- **Notebook:** TensorBoard_Profiling