


DL Optimizations for IoT Devices

Tolerance to Approximations in Deep Learning

Deep Learning Models Tolerate Approximations

- Most machine learning models, and deep learning ones in particular, are known to be quite *resilient* to various kinds of approximations.
- Changing slightly the **input data** or the **internal computations**, often doesn't change the final result (e.g. predicted class label).




Deep Learning Models Tolerate Approximations

- We can exploit this tolerance to approximations to make our models **faster, smaller and more efficient**
- Intuitively, we can approximate complex computations with simpler ones, large data with more compact representations, etc.
- This is an instance of a paradigm called **Approximate Computing**, which does not only apply to ML, but also to other domains, such as multimedia processing, optimization, etc.

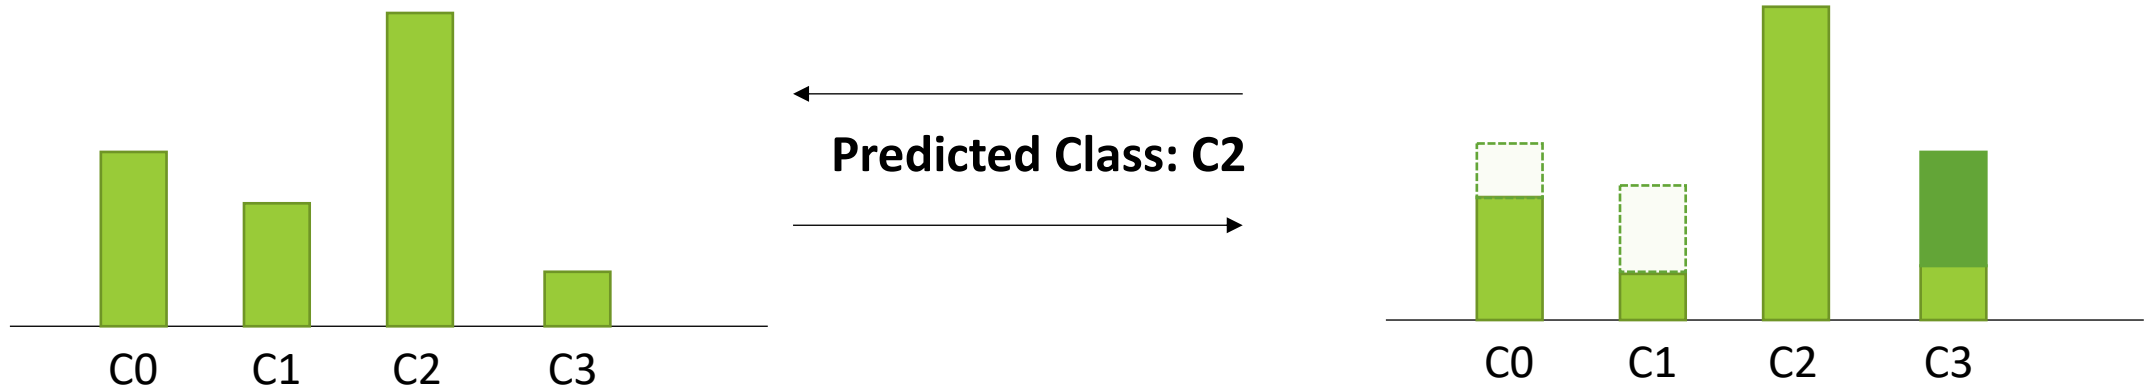
Deep Learning Models Tolerate Approximations

- Yes, but why are DL models tolerant to approximations?
- **Reason 1:** Gradient descent-based training algorithms converge even in presence of small “deviations” from the correct gradient direction of a mini-batch.



Deep Learning Models Tolerate Approximations

- Yes, but why are DL models tolerant to approximations?
- **Reason 2:** often we don't care about the exact output value. For instance, the output of a NN for multi-class classification are probability scores for each class. But finally, for most applications, we only care about the *argmax* of the output array, not about the exact probability values.




Deep Learning Models Tolerate Approximations

- Yes, but why are DL models tolerant to approximations?
- **Reason 3:** (linked with #2) DL models are highly redundant. They approximate a function using many more parameters and computations than those that are actually needed to reach that accuracy.
- Why (intuitively)?
 - If we knew the target function, we would be able to remove this redundancy, but we don't.
 - So, we normally “over-design” our model to make sure that it has a sufficient *capacity* to create a mapping that is close enough to the original function.

THE MOST IMPORTANT

Deep Learning Models Tolerate Approximations

- Because of “Reason 3”, in many cases, the **best way** to improve the efficiency of a DL model is to **simplify the network architecture!**
- Reduce # of layers, change type of layer (e.g. depth-wise/group-wise convolution), downscale input, etc.
- This often yields more efficient networks without accuracy loss!!
 - Example: SqueezeNet vs AlexNet (50x less parameters, same accuracy)



Deep Learning Models Tolerate Approximations

- However, architectural improvements still require human creativity
 - NN architecture design has replaced feature engineering in classic ML
 - Hard, time-consuming, manual process.
- Fortunately, we also have systematic ways to simplify Deep Learning models by inserting approximations. In particular, we'll focus on:
 - 1. Data Quantization**
 - 2. Network Pruning**
 - 3. Neural Architecture Search**
 - 4. Adaptive/dynamic models**

Deep Learning Models Tolerate Approximations

- Importantly, these techniques can yield performance/energy benefits only thanks to a **synergy between hardware and ML models**. In other words:
 - Model simplifications must be designed in such a way that they **actually result in a more efficient execution** on the target hardware...
 - ...or vice versa, the **hardware must be designed to exploit a given model simplification!**

Deep Learning Models Tolerate Approximations

- Importantly, these techniques can yield performance/energy benefits only thanks to a **synergy between hardware and ML models**. In other words:

- Model simplifications must be designed in such a way that they **actually result in a more efficient execution** on the target hardware...

- ...or vice versa, the hardware must be designed to exploit a given model simplification!

- Since this is not a HW design course, we'll focus mostly on model optimizations that can yield benefits on general purpose HW (MCUs, CPUs, etc.)
 - Particularly useful for edge devices, where the HW platform is typically general purpose