

# Machine Learning for IoT

## Lab 4 – Optimization

---

### **Exercise 1: Post-Training Quantization (PTQ)**

- 1.1. Modify the training scripts (see Lab 3) to save the Tensorflow datasets on disk. For instance, to save the temperature and humidity datasets, use the following commands:

```
tf.data.experimental.save(train_ds, './th_train')
tf.data.experimental.save(val_ds, './th_val')
tf.data.experimental.save(test_ds, './th_test')
```

You can load the dataset in other scripts with the following commands:

```
tensor_specs = (tf.TensorSpec([None, 6, 2], dtype=tf.float32),
                tf.TensorSpec([None, 2]))
train_ds = tf.data.experimental.load('./th_train', tensor_specs)
val_ds = tf.data.experimental.load('./th_val', tensor_specs)
test_ds = tf.data.experimental.load('./th_test', tensor_specs)
```

- 1.2. On your notebook, write a Python script to evaluate the prediction quality of TFLite models. Verify that the FP32 TFLite models have the same prediction quality as the Keras models.
- Write an evaluation script to compute the MAE of TFLite multi-output models, using *numpy* methods. Run this script to evaluate temperature and humidity forecasting models in TFLite format.
  - Write an evaluation script to compute the classification accuracy of TFLite models. Run this script to evaluate the keyword spotting models in TFLite format.

**N.B.:** Set the batch size to 1 when running inference with TFLite models:

```
test_ds = test_ds.unbatch().batch(1)
```

- 1.3. Apply Weights-Only PTQ to the multi-output models for temperature and humidity forecasting.

- Generate the TFLite models.

```
converter = tf.lite.TFLiteConverter.from_saved_model(saved_model_dir)
converter.optimizations = [tf.lite.Optimize.DEFAULT]
tflite_quant_model = converter.convert()
```

- Save the TFLite model on disk and measure the file size.
- Evaluate the MAE of each model using the script of 1.1.
- Plot TFLite Size vs. Temperature MAE and TFLite Size vs Humidity MAE for each model (MLP, CNN-1D) and precision (FP32 and INT8).
- Deploy the quantized models on the Raspberry and measure the inference latency. Plot Latency vs. Temperature MAE and Latency vs Humidity MAE.

- 1.4. Apply Weights-Only PTQ to the keyword spotting models. Plot TFLite Size vs Accuracy for the different pre-processing strategies and models.
- 1.5. Apply PTQ (Weights + Activations) to the multi-output models for temperature and humidity forecasting.

```
def representative_dataset_gen():
    for x, _ in train_ds.take(1000):
        yield [x]

converter = tf.lite.TFLiteConverter.from_saved_model(saved_model_dir)
converter.optimizations = [tf.lite.Optimize.DEFAULT]
converter.representative_dataset = representative_dataset_gen
tflite_quant_model = converter.convert()
```

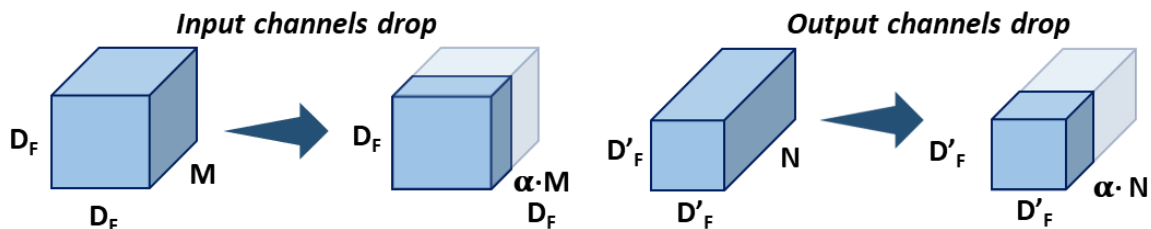
- Plot TFLite Size vs. Temperature MAE and TFLite Size vs Humidity MAE for each model and precision.
- Plot Latency vs Accuracy for each model and precision.

**N.B.:** PTQ (Weights + Activations) of CNN-1D is not supported.

- 1.6. Apply PTQ (Weights + Activations) to the keyword spotting models.
  - Plot TFLite Size vs Accuracy.
  - Plot Latency vs Accuracy.
- 1.7. For each task, which optimized model achieves the best accuracy-size-latency trade-off?

## **Exercise 2: Structured Pruning via Width Scaling**

- 2.1. Modify the models for temperature and humidity forecasting to implement models with a parametric *width*:
  - Multiply the number of filters (of convolutional layers) and the number of units (of dense and LSTM layers) by a parameter called the width multiplier  $\alpha \in (0, 1]$ . All the layers share the same width multiplier.
  - Train the models with different values of the width multiplier, i.e. 0.5 and 0.75. Evaluate the MAE, the TFLite Size, and Latency.



- 2.2. Repeat the same experiments with the models for keyword spotting.
- 2.3. Comment the results. Which is the best strategy for size optimization between quantization and structured pruning? Which is the best strategy for latency optimization?

## **Exercise 3: Magnitude-Based Pruning**

### 3.1. Train the multi-output models for temperature and humidity forecasting with magnitude-based pruning

- Define the sparsity scheduler as follows:

```
import tensorflow_model_optimization as tfmot

pruning_params = {'pruning_schedule':
                  tfmot.sparsity.keras.PolynomialDecay(
                      initial_sparsity=0.30,
                      final_sparsity=0.8,
                      begin_step=len(train_ds)*5,
                      end_step=len(train_ds)*15)
                  }

prune_low_magnitude = tfmot.sparsity.keras.prune_low_magnitude
model = prune_low_magnitude(model, **pruning_params)
```

- Define the pruning callback:

```
callbacks = [tfmot.sparsity.keras.UpdatePruningStep()]
```

- Train the model:

```
input_shape = [32, 6, 2]
model.build(input_shape)
model.fit(train_ds, epochs=20, validation_data=val_ds, callbacks=callbacks)
```

- Strip the model

```
model = tfmot.sparsity.keras.strip_pruning(model)
```

- Generate the TFLite model.
- Compress the model using the zlib and save the output on disk.

```
import zlib
tflite_model = converter.convert()
with open(tflite_model_dir, 'wb') as fp:
    tflite_compressed = zlib.compress(tflite_model)
    fp.write(tflite_compressed)
```

- Repeat the same experiment with *final\_sparsity=0.9*

### 3.2. Repeat the same experiments with the keyword spotting models.

### 3.3. Evaluate the prediction quality, the TFLite size, and the latency. Comment the results.